

Boundary Labeling for annotated documents

BACHELOR'S THESIS

submitted in partial fulfillment of the requirements for the degree of

Bachelor of Science

in

Bachelor's programme Software & Information Engineering

by

Jakob Klinger

Registration Number 1125755

to the Faculty of Informatics
at the TU Wien

Advisor: Ass.Prof. Dipl.-Inform. Dr.rer.nat Martin Nöllenburg

Assistance: Univ.Ass. Fabian Klute, M.Sc., B.Sc.

Vienna, 25th September, 2017

Jakob Klinger

Martin Nöllenburg

Erklärung zur Verfassung der Arbeit

Jakob Klinger
Scherzergasse 10/8, 1020 Wien

Hiermit erkläre ich, dass ich diese Arbeit selbständig verfasst habe, dass ich die verwendeten Quellen und Hilfsmittel vollständig angegeben habe und dass ich die Stellen der Arbeit – einschließlich Tabellen, Karten und Abbildungen –, die anderen Werken oder dem Internet im Wortlaut oder dem Sinn nach entnommen sind, auf jeden Fall unter Angabe der Quelle als Entlehnung kenntlich gemacht habe.

Wien, 25. September 2017

Jakob Klinger

Contents

Contents	v
1 Introduction	1
1.1 Terminology and Fundamentals	1
1.2 Related Work	3
2 The Algorithm	5
2.1 Problem Specification	5
2.2 Description	6
3 Implementation	9
3.1 Challenges	9
4 Evaluation and Testing	11
4.1 Data generation	11
4.2 Testing methods	11
4.3 Results	11
5 Conclusion	13
5.1 Further notes	13
Bibliography	15

Introduction

Whenever additional information needs to be inserted into an existing document without altering the original text, we can make use of annotations. They usually take the form of footnotes, which require only a minimal reference in the main text, and are used for a variety of reasons - for example, to provide additional information that would hinder the text's flow if inserted directly, or as a result of a commenting tool that is used for communication between an author and their editor. If a more obvious connection between the text and the referenced content is required, for example when lengthy comments are added, or if a change-tracking tool is used, the reference is often placed to the side of the document and visibly connected to the part of the text it is referring to. This style of annotation is easily implemented on virtual documents, since they can be hidden on demand, however if the annotations need to be included in a printed version, there are several issues that arise regarding readability of the final product and ambiguity of text-annotation assignments.

In this thesis, we will look at ways to use Boundary Labeling for this problem, which means that all annotations will be placed somewhere outside of the text they are referencing and will be visually connected to the feature they are referencing. (See also [3])

The guidelines on how to create suitable labelings are as follows: the connections should be as direct as possible, no important information should be obscured, and it should be easily discernable which Label belongs where. These three criteria easily come into conflict with one another, as the text usually is very dense and leaves little space for lines in between, yet one shouldn't allow them to pass through the text, as this makes the text harder to read.

1.1 Terminology and Fundamentals

Boundary Labeling (or equivalent concepts) can be applied to a space with different geometry or more dimensions, but this thesis will only concern itself with two-dimensional,

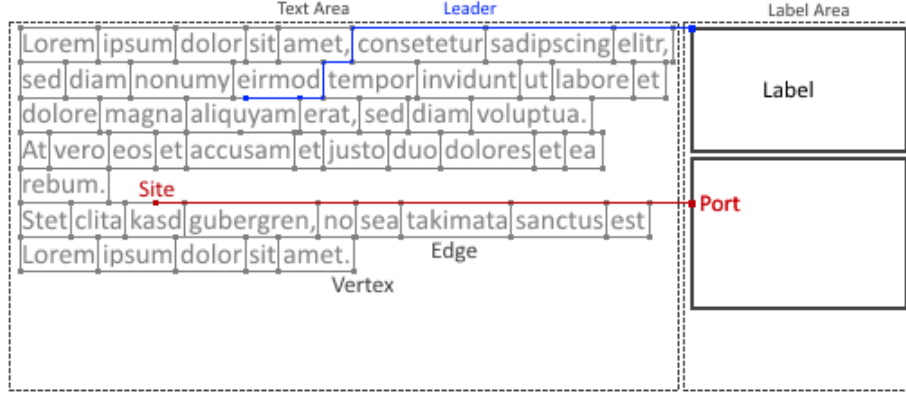


Figure 1.1: Illustrated guide to the labeling terminology

Euclidean space. To easily reference important concepts, some additional terminology will be introduced as well. (See Fig. 1.1 for a visual explanation)

A *graph* $G = \langle V, E \rangle$ is a tuple of *vertices* $V = \{v_1, v_2, \dots, v_n\}$ and *edges* $E = \{e_1, e_2, \dots, e_m\}$. A vertex v is a featureless object. Each edge e is a relation between two vertices $E \subseteq V \times V$. We call two vertices $u, v \in V$ *adjacent*, if the edge $e = (u, v) \in E$. A *path* $P = v_1, \dots, v_h$ is an ordered sequence of vertices, where each vertex must have an edge connecting it to the subesquent one. *Depth-first search* is a searching algorithm on a graph G that starts at a given vertex $v \in V$ and explores the graph by traversing its edges as far as possible before backtracking, and continues to do so until a pre-defined goal is met.

Polylines are a connected series of *line segments*. Line segments are straight lines that contain each point between their starting and end point. *Labels* hold additional information and are represented as boxes containing this information. They are usually placed in the *label area* which is a rectangular area designated to hold labels. It is located next to of the bigger, rectangular *text area*, which contains the document's text and all *sites*, the points or objects that a label's information refers to. **If multiple label areas exist on different sides of the text area, we speak of *multi-sided labeling*, otherwise we speak of *one-sided labeling*. We will be using one-sided labeling in our implementation. Some space was left in between the text and label area, to make connecting sites to their labels easier, which we will call the *routing area*.** The site and the label are connected via a *leader*, a polyline that can be further classified by looking at the orientation of its segments: *O-Segments* run orthogonally to the border of the label area. *P-Segments* run parallel to the border of the label area, and as such must be combined with other segments for the leader to reach its destination. *S-Segments* are not required to have any particular orientation, and simply connect their start and ending points in a straight line. The leader's name is created by combining the name of the segments - for example, the blue leader from Fig 1.1 would be classified as an OPOPO-Leader. The location where a leader connects to the label is called the *port*. It can be restricted to pre-defined positions.

1.2 Related Work

Boundary labeling was first introduced by Bekos et al. in 2004 (see [3]), where both one-sided and multi-sided labelings with different leader types are looked into. They also showed that the optimal placement of arbitrarily-sized labels on two sides of the text area can be NP-hard by drawing comparisons to the Partition-Problem. However, a pseudo-polynomial solution exists for this problem, which was adapted to this variation of the problem.

Since then, several papers have been written about boundary labeling. One of these is [2], which looks into the readability of different leader styles. Interestingly, some leader styles perform quite well, despite the study's participants preferring others over them, with OPO-Leaders being both least preferred and the hardest to follow.

Another article using boundary labeling is [4] by Götzelmann et al., which creates boundary labeling-style annotations along other methods to label different parts of three-dimensional figures, resulting in pictures similar to what could be found in a textbook. As this algorithm works in real-time, it is suitable for labeling interactive models and allows for user interaction.

Boundary labeling in text documents however, is rarely discussed, and only few papers exist about this topic. The programs that employ this style of annotation often also use rather simple algorithms, to mediocre results or make extensive use of the interactivity of a digital medium, showing annotations only on demand. However, the few papers that approach this topic add interesting information to the discussion.

The paper about the *Luatodonotes-Package*[5] uses several styles of drawing leaders, and came to the conclusion that leaders without bends are easier to follow, which fits with [2]'s observations, which ranks OPO- and PO-Leaders lower than other variants. However, most solutions proposed in [5] do not consider whether a path overlaps with text or not, which results in a decrease in readability. While we do not use the routing and leader styles introduced in this paper, the results can be used in comparisons regarding readability of the main text and ease of use of the different leader styles.

The thesis by Loose[7] on the other hand is based around only using the free space between lines and words, which produces longer leaders, and forces curves, but doesn't obscure any part of the text. The two different approaches in this paper were a clustering-based algorithm, which was previously described in [8], and a flow network-based approach. Several concepts of this paper, such as the graph-based strategy and the usage of a routing area will be adopted in our thesis and it is by far the biggest influence on our approach to the problem.

Lin et al.[6] use only OPO-Leaders that have their P-Segment located outside of the text area in their paper, but allow the leaders to use the text area's border on the opposite side of the label area to route upwards or down. This allows for more labels to be placed as close as possible to their leader's source, at the cost of increasing select leaders' length and placing some labels out of order. While this is an interesting way to avoid longer leaders in general, it is hard to combine with the graph-based routing that happens inside the text area, so we won't make use of it.

The Algorithm

2.1 Problem Specification

In our algorithm, the focus was put on keeping the text as readable as possible. This means that leaders aren't allowed to pass through words - they can only use the space between words and in between lines. Additionally, we wanted to maximize the usage of the label area's space, even if we have to move labels further from their sites than needed. Finally, we wanted to keep the leader's length low for any given combination of labels and site placements, so every segment of a leader has to decrease the remaining distance between the site and its label's port.

These properties were modelled by creating a rectangle around each word that no leaders may pass through. If this is done for all words in a paragraph, the remaining available space for routing resembles a grid of perfectly parallel streets. **NOTE: Insert example image for this and graph.** Since the layout encourages the usage of P- and O-segments, we decided to represent the possible routes by creating a graph which has its vertices located whenever multiple directions are available - in between lines, above or below whitespace characters. The graph's edges represent the different possible segments of a leader and connect each vertex with the closest vertices directly above, below or to the sides, given that no text is blocking the connection. **The resulting graph looks similar to the one in Fig. <TODO>**

For ease of visualisation, we limited the number of leaders per edge to one, and forbid the leaders to cross one another, as this could be mistaken for two curves touching each other at the corners. We also decided to place the annotations in the order in which they appear in the text, and always place their port above the leader's source, which means that crossings are inevitable if a leader is located above and to the left of an earlier leader. To ensure this never happens, leaders aren't allowed to enter the top left quadrant of a source vertex, unless no routing could be found for that vertex, which means there won't be any leaders to collide with. Additionally, to keep the leader's length minimal, each leader is only allowed to move towards the source, which places

a hard limit on how far up a leader is allowed to go, and restricts the possible choices at each node to edges that allow the leader to move in that direction. To maximize the space available to subsequent annotations, each label is placed as far up as possible. This also means that if there's no placement available that is above its leader's source, the annotation cannot be placed, as we previously restricted the leader's movement.

2.2 Description

```

Data: A text with annotations, stored as a String-Array
Result: A Graph (as described above)
1 initialization
2 foreach w in words do
3   if w is annotation then
4     v ← new Vertex(previousWord.getCenter())
5     v.setAnnotation(w)
6     Graph.addVertex(v)
7     UpperVerticesList.addVertex(v)
8   else
9     if w is too big for the line then
10      startNewLine()
11      connectBasedOnPosition(UpperVerticesList)
12      UpperVerticesList ← LowerVerticesList
13      emptyList(LowerVerticesList)
14    end
15    v1 ← new Vertex(w.getTopLeft())
16    v2 ← new Vertex(w.getTopRight())
17    v3 ← new Vertex(w.getBottomLeft())
18    v4 ← new Vertex(w.getBottomRight())
19
20    Graph.addAll(v1, v2, v3, v4)
21    UpperVerticesList.addAll(v1, v2)
22    LowerVerticesList.addAll(v3, v4)
23    Graph.createEdgeBetween(v1, v3)
24    Graph.createEdgeBetween(v2, v4)
25  end
26 end

```

Algorithm 1: Representation of the Graph-creation algorithm in pseudocode

To create leaders that exclusively use the space not taken up by a word's bounding rectangle, we decided to use a graph similar to the one Loose [7] used. As each vertex represents a physical location, they will have co-ordinates associated with them, and the vertices representing the sites will hold additional information regarding its leader and

label. The graph is constructed by placing vertices between the lines, located next to each corner of a word's bounding rectangle, with two consecutive words in a line sharing the two vertices associated with their adjacent corners. For the sites, we inserted an additional vertex above the center of the word, which will serve as the leader's starting point. After placing all vertices, we created edges between each vertex and the closest horizontal neighbour to both sides, and between nodes that are located exactly above or below each other, and exactly one line apart. (For a representation in Pseudocode, see Alg. 1)


In our labeling algorithm, we decided to work through the labels in the order they appear in the text, placing each as far up as possible, and skipping any label that would've required us place the label below its leader's source node. We also opted to use fixed ports located in the top left corner of each label, as this allowed us to unambiguously place annotations by only knowing their port location, which is equal to their leader's ending point. To not restrict the label's placement by the line spacing, we left a buffer zone between the text area and the label area, which allows us to place labels even further above to optimally use the available space. The algorithm's input consists of the Graph G and a set of sites $V_{ann} \subset V$ which will be routed using a depth-first search algorithm, restricted to only selecting vertices located above or to the right of the current vertex, prioritizing moving up whenever possible. After the algorithm terminates, it returns a set with routing information for each vertex, which contains a Path $P = \{v_1, \dots, v_n\}$, leading from the site to the text area's border, along with some additional information on how to draw the OPO-Segment that connects to the Label's port. If the routing for any given site failed, the path consists of a single vertex - the site. (For an illustration in Pseudocode see Alg. 2.)

```

Data: A single annotation's source and its Graph
Result: A List of vertices describing the Leader's Path
1 initialization
2 while currentVertex not at right text border do
3   if (Graph.getTopNeighbourOf (currentVertex)  $\neq$  null)  $\wedge$   $\neg$ backtracking
4     then
5       Path.addVertex (currentVertex)
6       currentVertex  $\leftarrow$  Graph.getTopNeighbourOf (currentVertex)
7     else if Graph.getRightNeighbourOf (currentVertex)  $\neq$  null then
8       Path.addVertex (currentVertex)
9       currentVertex  $\leftarrow$  Graph.getTopNeighbourOf (currentVertex)
10      backtracking  $\leftarrow$  False
11    else
12      backtracking  $\leftarrow$  True
13      repeat
14        oldVertex  $\leftarrow$  currentVertex
15        currentVertex  $\leftarrow$  Path.getLastEntry ()
16        Path.RemoveVertex (currentVertex)
17      until currentVertex's Position is below oldVertex or Path is Empty
18
19      if currentVertex not below oldVertex then //No path found
20        break
21      end
22    end
23 end

```

Algorithm 2: The Depth-First-Search algorithm used in the program.



Implementation

The program was written in Java 1.8.0u40, using JGraphT1.0.1[1] as graph library. Since we only want to create leaders that don't intersect with the text, the graph was created alongside the placement of the words on the canvas.

3.1 Challenges

CHAPTER 4

Evaluation and Testing

- 4.1 Data generation
- 4.2 Testing methods
- 4.3 Results

CHAPTER 5



Conclusion

5.1 Further notes

Bibliography

- [1] Jgrapht - a free java graph library.
- [2] Lukas Barth, Andreas Gemsa, Benjamin Niedermann, and Martin Nöllenburg. *On the Readability of Boundary Labeling*, pages 515–527. Springer International Publishing, Cham, 2015.
- [3] Michael A. Bekos, Michael Kaufmann, Antonios Symvonis, and Alexander Wolff. *Boundary Labeling: Models and Efficient Algorithms for Rectangular Maps*, pages 49–59. Springer Berlin Heidelberg, Berlin, Heidelberg, 2005.
- [4] Timo Götzelmann, Knut Hartmann, and Thomas Strothotte. *Agent-Based Annotation of Interactive 3D Visualizations*, pages 24–35. Springer Berlin Heidelberg, Berlin, Heidelberg, 2006.
- [5] Philipp Kindermann, Fabian Lipp, and Alexander Wolff. *Luatodonotes: Boundary Labeling for Annotations in Texts*, pages 76–88. Springer Berlin Heidelberg, Berlin, Heidelberg, 2014.
- [6] C. C. Lin, H. Y. Wu, and H. C. Yen. Boundary labeling in text annotation. In *2009 13th International Conference Information Visualisation*, pages 110–115, July 2009.
- [7] Amrei Loose. Annotation von texten unter berücksichtigung von textzwischenräumen. Master’s thesis, Karlsruhe Institute of Technology, 2015.
- [8] Martin Nöllenburg, Valentin Polishchuk, and Mikko Sysikaski. Dynamic one-sided boundary labeling. In *Proceedings of the 18th SIGSPATIAL International Conference on Advances in Geographic Information Systems*, GIS ’10, pages 310–319, New York, NY, USA, 2010. ACM.